

Differentially Private Algorithm and Auction Configuration

Ellen Vitercik

CMU, Theory Lunch

October 11, 2017

Joint work with Nina Balcan and Travis Dick





Prices learned from purchase histories can reveal information about individual purchases.



We need a way to **privately** set prices and design auctions based on purchase histories.









Dear Lab
Technician,

Please update your
default CPLEX
parameters to ...



Suppose a parameter correlates with a certain disease.

An attacker can infer information about medical records used to tune those parameters.



We need a way to **privately** configure algorithms.

Many works have shown that it is possible to invert a machine learning model to **infer sensitive information** about its training set.



By observing a series of recommendations from websites such as Amazon, an adversary can **infer individual users' purchases.** [Calandrino et al. 2011]

Recommendations for you in Grocery & Gourmet Food




It is possible to **extract images of training subjects** from facial recognition models. The attacker has only the person's name and access to a facial recognition system that returns a class confidence score. [Fredrikson et al. 2015]




It is possible to invert a machine learning model to learn **sensitive genomic information** about individuals.


[Fredrikson et al. 2014]




In response, computer scientists have developed private machine learning algorithms.



Existing private ML algorithms apply to optimization problems defined by well-studied, well-understood functions.




What if the objective is nonconvex and not differentiable?



We provide a **private** algorithm for maximizing data-dependent piecewise Lipschitz functions.

Applications in:



Algorithm configuration and mechanism design reduce to maximizing data-dependent piecewise Lipschitz functions.

Introduction

Setup

- ➡ Overview: Algorithm configuration and auction design
 - Differential privacy
 - Private pricing design

The algorithm

- Privacy guarantees


- Utility guarantees

- Example: private multi-item pricing

Summary

Learning-based algorithm configuration:


Tune algorithm parameters to achieve high performance over a specific application domain.



Led to breakthroughs in:

- Combinatorial auctions
[Leyton-Brown et al., 2009]
- Scientific computing
[Demmel et al., 2005]
- Vehicle routing
[Caseau et al., 1999]
- SAT
[Xu et al., 2008]

Learning-based algorithm configuration




Algorithm Designer

Algorithm

How can I use the set of samples to find an algorithm that's best for my application domain?

We show that our private algorithm has strong utility guarantees for many **algorithm configuration problems**.


Greedy algorithm configuration



Hard combinatorial problems show up in diverse domains where **privacy preservation is crucial**.

We show that our private algorithm has strong utility guarantees for many **algorithm configuration problems**.

Greedy algorithm configuration




These are often solved by greedy algorithms where elements are iteratively added to a solution set according to a **heuristic**.

E.g., in knapsack:

$$\frac{(\text{size of item } i)}{(\text{value of item } i)}$$

We show that our private algorithm has strong utility guarantees for many **algorithm configuration problems**.

Greedy algorithm configuration




Gupta and Roughgarden [2017] proposed an **infinite family of greedy heuristics** for the knapsack and max weight independent set problems.

E.g., in knapsack:

$$\frac{(\text{size of item } i)}{(\text{value of item } i)^\rho}$$

We show that our private algorithm has strong utility guarantees for many **algorithm configuration problems**.

Greedy algorithm configuration




Our private algorithm uses sample problem instances to find a **nearly optimal greedy heuristic** for the specific application domain.

No **sensitive information** about the training set is revealed.

We show that our private algorithm has strong utility guarantees for many **algorithm configuration problems**.


Integer quadratic programming algorithm configuration



IQPs are used in many applications where privacy preservation is essential, such as financial portfolio optimization.

We show that our private algorithm has strong utility guarantees for many **algorithm configuration problems**.

Integer quadratic programming algorithm configuration




IQPs are often approximated by solving a semi-definite program and **rounding** the vectors to integer values.

There are many different rounding schemes with varying quality.

We show that our private algorithm has strong utility guarantees for many **algorithm configuration problems**.

Integer quadratic programming algorithm configuration



Our private algorithm uses sample IQP instances to find a **nearly optimal rounding scheme** for the specific application domain.

No **sensitive information** about the training set is revealed.

Learning-based mechanism design:

Use information about past consumers to design mechanisms that extract high revenue from future consumers.

Employed throughout industry.



Garnered significant attention in TCS.
[Elkind, 2007, Cole and Roughgarden, 2014, Huang et al., 2015, Medina and Mohri, 2014, Morgenstern and Roughgarden, 2015, Devanur et al., 2016, etc.]

Introduction

Setup

Overview: Algorithm configuration and auction design



Differential privacy

Private pricing design

The algorithm

Privacy guarantees

Utility guarantees

Example: private multi-item pricing

Summary






Differential privacy




An algorithm, given an input dataset D , is **differentially private** if the following holds:

The output reveals (almost) nothing more about a record in D than the output would have if the record wasn't contained in D .





Algorithm




An algorithm \mathcal{A} is (ϵ, δ) -**differentially private** if for all pairs of neighboring datasets D, D' and all sets \mathcal{O} of outputs,

$$\mathbb{P}[\mathcal{A}(D) \in \mathcal{O}] \leq e^\epsilon \mathbb{P}[\mathcal{A}(D') \in \mathcal{O}] + \delta$$

↑

$e^\epsilon \approx 1 + \epsilon$



Introduction

Setup

Overview: Algorithm configuration and auction design

Differential privacy

➡ Private pricing design

The algorithm

Privacy guarantees

Utility guarantees

Example: private multi-item pricing

Summary

Single-item pricing problem:

One good for sale 

Single-item pricing problem:

Distribution over buyers:  $\sim \mathcal{D}$

Buyers' values are denoted as: $\text{value}\left(\text{img alt="green person icon" data-bbox="671 261 691 281"}, \text{img alt="coffee cup icon" data-bbox="711 258 788 338"}\right)$

$\text{Revenue}\left(\text{img alt="orange person icon" data-bbox="294 396 314 416"}, \rho\right) = \rho \cdot \mathbf{1}\left[\text{value}\left(\text{img alt="orange person icon" data-bbox="576 396 596 416"}, \text{img alt="coffee cup icon" data-bbox="616 393 693 473"}\right) \geq \rho\right]$


Pricing algorithm receives a set $\mathcal{S} = \left\{ \text{img alt="blue person icon" data-bbox="658 538 678 558"}, \text{img alt="orange person icon" data-bbox="691 538 711 558"}, \dots, \text{img alt="yellow person icon" data-bbox="771 538 791 558"} \right\} \sim \mathcal{D}^N$

Goal: maximize $\frac{1}{N} \left(\text{Revenue}\left(\text{img alt="blue person icon" data-bbox="576 648 596 668"}, \rho\right) + \dots + \text{Revenue}\left(\text{img alt="yellow person icon" data-bbox="868 648 888 668"}, \rho\right) \right)$

Learning theory tells us that this value of ρ
approximately maximizes $\mathbb{E}\left[\text{Revenue}\left(\text{img alt="black person icon" data-bbox="788 821 808 841"}, \rho\right)\right]$

$$\mathcal{S} = \{ \text{person} \} \sim \mathcal{D}$$

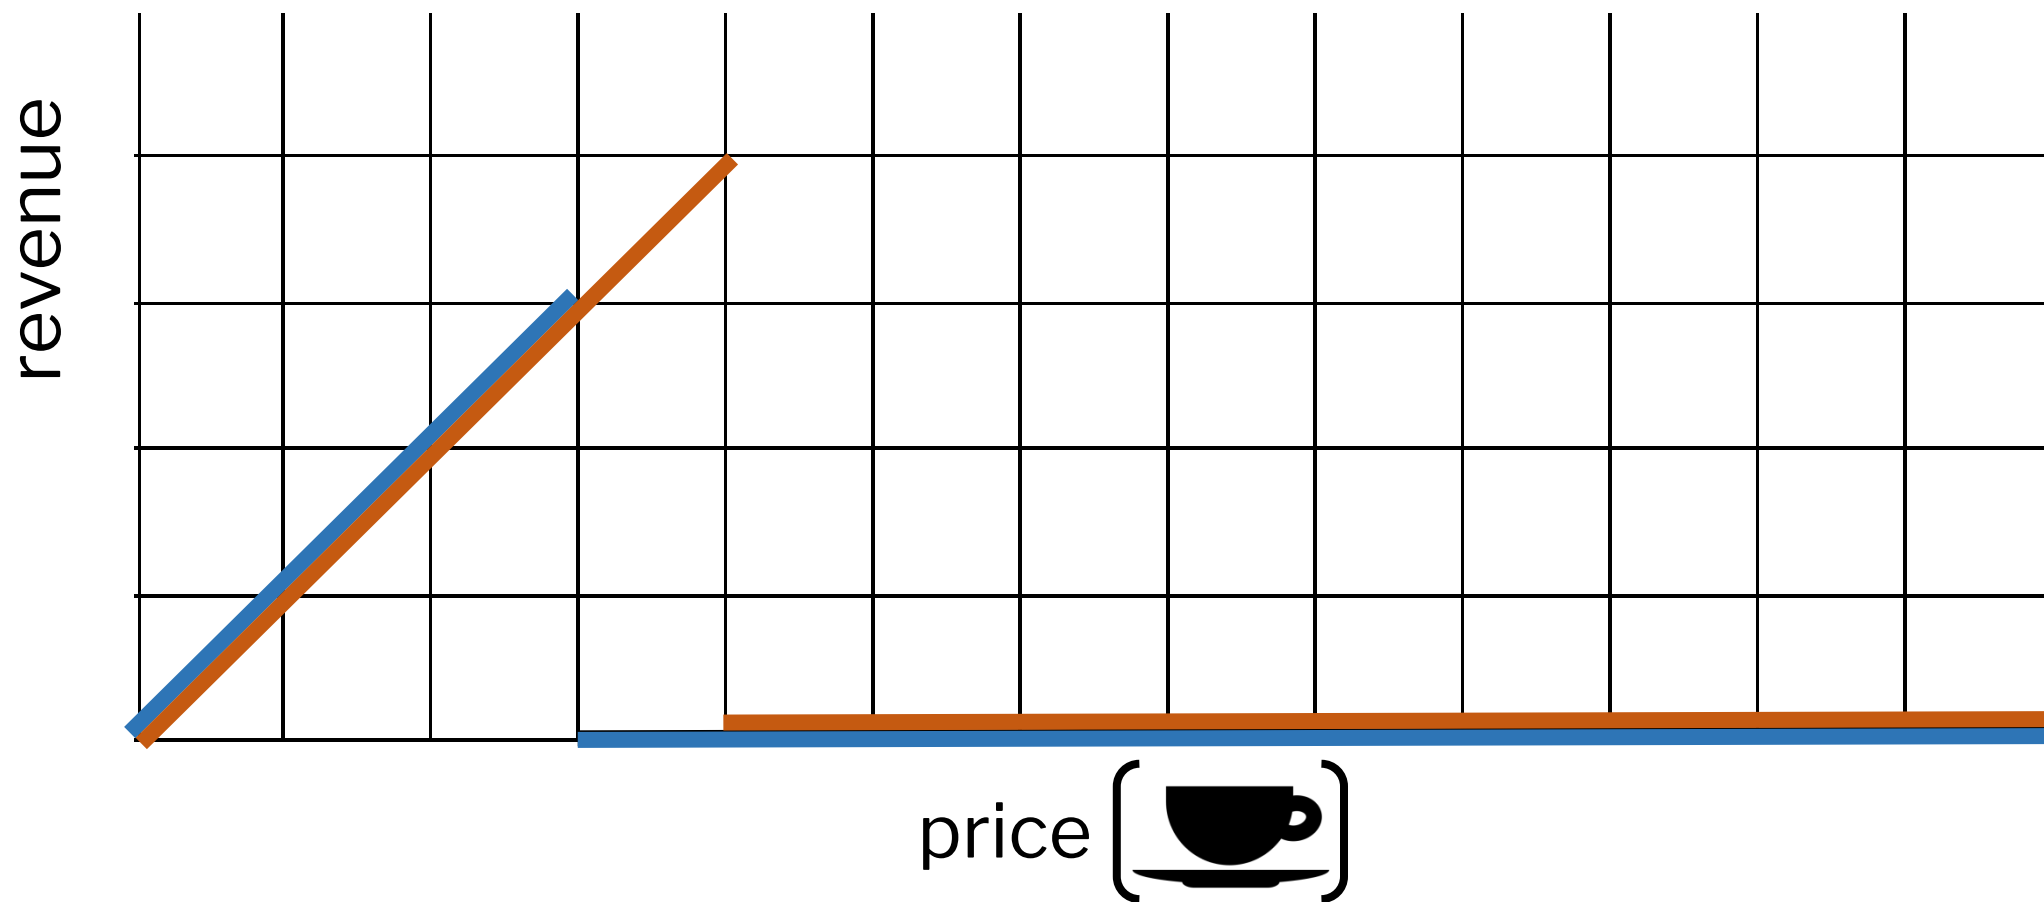
$$\text{value}(\text{person}, \text{cup}) = 3$$



$$\mathcal{S} = \{ \text{blue person}, \text{orange person} \} \sim \mathcal{D}^2$$

$$\text{value} \left(\text{blue person}, \text{cup} \right) = 3$$


$$\text{value} \left(\text{orange person}, \text{cup} \right) = 4$$




$$\mathcal{S} = \{ \text{blue person}, \text{orange person} \} \sim \mathcal{D}^2$$

$$\text{value} \left(\text{blue person}, \text{cup} \right) = 3$$


$$\text{value} \left(\text{orange person}, \text{cup} \right) = 4$$



We want to write an algorithm that gets average revenue as close to \$3 as possible while preserving differential privacy.




We want to write an algorithm that gets average revenue as close to \$3 as possible while preserving differential privacy.




General problem:

Given a piecewise utility function $U_{\mathcal{S}}(\rho)$, privately find a parameter $\hat{\rho}$ that approximately maximizes $U_{\mathcal{S}}(\rho)$.



General problem:

Given a piecewise utility function $U_{\mathcal{S}}(\rho)$, privately find a parameter $\hat{\rho}$ that approximately maximizes $U_{\mathcal{S}}(\rho)$.



Introduction

Setup

Overview: Algorithm configuration and auction design

Differential privacy

Private pricing design

➡ The algorithm

Privacy guarantees

Utility guarantees


Example: private multi-item pricing

Summary


Ideally, we want to use the exponential mechanism
[McSherry and Talwar 2007]

Pick a parameter vector $\boldsymbol{\rho}$ with probability proportional to
$$\exp(c \cdot \varepsilon \cdot U_{\mathcal{S}}(\boldsymbol{\rho}))$$


Sampling from $\exp(c \cdot \varepsilon \cdot U_{\mathcal{S}}(\rho))$ is easy when we can calculate its antiderivative and the parameter space is one-dimensional.



What about when the parameter space is multi-dimensional?




Generally, sampling from a multi-dimensional distribution is hard.



Existing efficient sampling techniques work for **logconcave distributions**.

[Lovász and Vempala 2006, 2007, Bassily et al. 2014]


When $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise concave, $\exp(c \cdot \varepsilon \cdot U_{\mathcal{S}}(\boldsymbol{\rho}))$ is logconcave on each piece.



Theorem (part 1)

Suppose $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz and concave over convex regions in \mathbb{R}^d .


There is an (ε, δ) -DP algorithm that approximately maximizes $U_{\mathcal{S}}(\boldsymbol{\rho})$. Its running time polynomial in $d, L, \varepsilon^{-1}, \log \delta^{-1}$, and $|\mathcal{S}|$.



Step 1:


Estimate the probability mass that the exponential mechanism places on each piecewise portion of the domain.

This requires an integral estimation algorithm. [Lovász and Vempala '07]



Step 2:

Sample a portion of the domain according to that probability.




Step 3:

Sample a point in that portion of the domain with probability approximately according to the exponential mechanism:

$$\exp(c \cdot \varepsilon \cdot U_{\mathcal{S}}(\boldsymbol{\rho}))$$

This requires approximate sampling [Bassily et al. 2014]



Proof sketch: Show that the output distribution is **almost** $\exp(c \cdot \varepsilon \cdot U_{\mathcal{S}}(\boldsymbol{\rho}))$.

Lemma [Bassily et al. 2014]

If an algorithm \mathcal{A} 's output distribution is $\tilde{\varepsilon}$ -close to that of a $(\varepsilon, 0)$ -differentially private algorithm for every input dataset \mathcal{S} , then \mathcal{A} is $(2\tilde{\varepsilon} + \varepsilon, 0)$ -differentially private.


(Close under the metric $D_{\infty}(\mu_f, \mu_g) = \sup |\log f(x)/g(x)|$.)

Since the exponential mechanism is $(\varepsilon, 0)$ -differentially private, our algorithm is $(O(\varepsilon), \delta)$ -differentially private.

Theorem (part 1)

Suppose $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz and concave over convex regions in \mathbb{R}^d .


There is an (ε, δ) -DP algorithm that approximately maximizes $U_{\mathcal{S}}(\boldsymbol{\rho})$. Its running time polynomial in $d, L, \varepsilon^{-1}, \log \delta^{-1}$, and $|\mathcal{S}|$.



Theorem (part 1)

Suppose $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz and concave over convex regions in \mathbb{R}^d .

There is an (ε, δ) -DP algorithm that **approximately** maximizes $U_{\mathcal{S}}(\boldsymbol{\rho})$. Its running time polynomial in $d, L, \varepsilon^{-1}, \log \delta^{-1}$, and $|\mathcal{S}|$.



Introduction

Setup

- Overview: Algorithm configuration and auction design

- Differential privacy

- Private pricing design

The algorithm

- Privacy guarantees

- ➔ Utility guarantees

 - Example: private multi-item pricing

Summary

Theorem (part 2)

Suppose $U_{\mathcal{S}}(\boldsymbol{\rho})$ is (w, k) -**dispersed**, piecewise L -Lipschitz, and the parameter space is contained in a ball of radius R .


Let $\hat{\boldsymbol{\rho}} \in \mathbb{R}^d$ be the parameter vector output by our (ε, δ) -DP algorithm.

With probability at least $1 - \gamma$,


$$U_{\mathcal{S}}(\hat{\boldsymbol{\rho}}) \geq \max U_{\mathcal{S}}(\boldsymbol{\rho}) - O\left(\frac{H}{|\mathcal{S}|\varepsilon} \left(d \log \frac{R}{w} + \log \frac{1}{\gamma}\right) + Lw + \frac{Hk}{|\mathcal{S}|}\right)$$

($U_{\mathcal{S}}(\boldsymbol{\rho})$ is an average of functions with range in $[0, H]$.)


Every private algorithm needs to add some noise to the parameter it outputs.




Therefore, nearby parameters should have similar utilities.




We often guarantee that nearby parameters have similar utilities by assuming the utility function is “nice.”



What does it mean for a piecewise function to be “nice”?




What does it mean for a piecewise function to be “nice”?




Let $\mathcal{P} = \{P_1, \dots, P_t\}$ be a partition of the parameter space such that $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz within each P_i .

$U_{\mathcal{S}}$ is **(w, k) -dispersed** if for any ball \mathcal{B} of radius w , the number of sets in \mathcal{P} that intersect \mathcal{B} is at most k .




Let $\mathcal{P} = \{P_1, \dots, P_t\}$ be a partition of the parameter space such that $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz within each P_i .

$U_{\mathcal{S}}$ is **(w, k) -dispersed** if for any ball \mathcal{B} of radius w , the number of sets in \mathcal{P} that intersect \mathcal{B} is at most k .



Let $\mathcal{P} = \{P_1, \dots, P_t\}$ be a partition of the parameter space such that $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz within each P_i .

$U_{\mathcal{S}}$ is **(w, k) -dispersed** if for any ball \mathcal{B} of radius w , the number of sets in \mathcal{P} that intersect \mathcal{B} is at most k .



Theorem (part 2)

Suppose $U_{\mathcal{S}}(\boldsymbol{\rho})$ is (\mathbf{w}, \mathbf{k}) -dispersed, piecewise L -Lipschitz, and the parameter space is contained in a ball of radius R .

Let $\hat{\boldsymbol{\rho}} \in \mathbb{R}^d$ be the parameter vector output by our (ε, δ) -DP algorithm.


With probability at least $1 - \gamma$,

$$U_{\mathcal{S}}(\hat{\boldsymbol{\rho}}) \geq \max U_{\mathcal{S}}(\boldsymbol{\rho}) - O\left(\frac{H}{|\mathcal{S}|\varepsilon} \left(d \log \frac{R}{\mathbf{w}} + \log \frac{1}{\gamma}\right) + L\mathbf{w} + \frac{H\mathbf{k}}{|\mathcal{S}|}\right)$$

($U_{\mathcal{S}}(\boldsymbol{\rho})$ is an average of functions with range in $[0, H]$.)


Let $\mathcal{P} = \{P_1, \dots, P_t\}$ be a partition of the parameter space such that $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz within each P_i .

$U_{\mathcal{S}}$ is **(w, k) -dispersed** if for any ball \mathcal{B} of radius w , the number of sets in \mathcal{P} that intersect \mathcal{B} is at most k .



Let $\mathcal{P} = \{P_1, \dots, P_t\}$ be a partition of the parameter space such that $U_{\mathcal{S}}(\boldsymbol{\rho})$ is piecewise L -Lipschitz within each P_i .

$U_{\mathcal{S}}$ is **(w, k) -dispersed** if for any ball \mathcal{B} of radius w , the number of sets in \mathcal{P} that intersect \mathcal{B} is at most k .



Theorem (part 2)

Suppose $U_{\mathcal{S}}(\boldsymbol{\rho})$ is (\mathbf{w}, \mathbf{k}) -dispersed, piecewise L -Lipschitz, and the parameter space is contained in a ball of radius R .

Let $\hat{\boldsymbol{\rho}} \in \mathbb{R}^d$ be the parameter vector output by our (ε, δ) -DP algorithm.

With probability at least $1 - \gamma$,

$$U_{\mathcal{S}}(\hat{\boldsymbol{\rho}}) \geq \max U_{\mathcal{S}}(\boldsymbol{\rho}) - O\left(\frac{H}{|\mathcal{S}|\varepsilon} \left(d \log \frac{R}{\mathbf{w}} + \log \frac{1}{\gamma}\right) + L\mathbf{w} + \frac{H\mathbf{k}}{|\mathcal{S}|}\right)$$

($U_{\mathcal{S}}(\boldsymbol{\rho})$ is an average of functions with range in $[0, H]$.)

Introduction

Setup

- Overview: Algorithm configuration and auction design

- Differential privacy

- Private pricing design

The algorithm

- Privacy guarantees

- Utility guarantees

- ➔ Example: private multi-item pricing

Summary

Multi-item pricing problem:

Multiple goods for sale



Multi-item pricing problem:

Distribution over buyers:  $\sim \mathcal{D}$

Buyers' values are denoted as: $\text{value}\left(\text{img alt="green person icon" data-bbox="635 245 655 325"}, \text{img alt="coffee cup icon" data-bbox="675 245 755 325"}\right)$

$\text{value}\left(\text{img alt="green person icon" data-bbox="635 355 655 435"}, \text{img alt="cupcake icon" data-bbox="685 355 735 435"}\right)$


$$\begin{aligned} \text{Revenue}\left(\text{img alt="orange person icon" data-bbox="255 465 275 545"}, \boldsymbol{\rho}\right) &= \rho_1 \cdot \mathbf{1}\left[\text{value}\left(\text{img alt="orange person icon" data-bbox="540 475 560 555"}, \text{img alt="coffee cup icon" data-bbox="580 475 660 555"}\right) \geq \rho_1\right] \\ &\quad + \rho_2 \cdot \mathbf{1}\left[\text{value}\left(\text{img alt="orange person icon" data-bbox="540 575 560 655"}, \text{img alt="cupcake icon" data-bbox="595 575 645 655"}\right) \geq \rho_2\right] \end{aligned}$$

Pricing algorithm receives a set $\mathcal{S} = \left\{ \text{img alt="blue person icon" data-bbox="620 700 640 780"}, \text{img alt="orange person icon" data-bbox="650 700 670 780"}, \dots, \text{img alt="yellow person icon" data-bbox="730 700 750 780"} \right\} \sim \mathcal{D}^N$

Goal: maximize $\frac{1}{N} \left(\text{Revenue}\left(\text{img alt="blue person icon" data-bbox="535 815 555 895"}, \boldsymbol{\rho}\right) + \dots + \text{Revenue}\left(\text{img alt="yellow person icon" data-bbox="825 815 845 895"}, \boldsymbol{\rho}\right) \right)$

The average revenue over $\mathcal{S} = \{\text{blue person}, \text{orange person}, \dots, \text{yellow person}\}$ of a pricing mechanism is piecewise linear and 1-Lipschitz.

$U_{\mathcal{S}}(\boldsymbol{\rho}) = \text{average revenue}$




How dispersed is $U_{\mathcal{S}}(\boldsymbol{\rho})$?

Let $f_i(v)$ be the density of the distribution over the “typical” buyer’s value for item i .

Suppose there are m items and $\max_{i,v} f_i(v) \leq \kappa$.

With probability at least $1 - \gamma$, $U_{\mathcal{S}}(\boldsymbol{\rho})$ is (w, k) -dispersed with $w = \frac{\gamma}{|\mathcal{S}| \kappa \sqrt{m}}$ and $k = m$.

$U_{\mathcal{S}}(\boldsymbol{\rho})$ = average revenue



Let $f_i(v)$ be the density of the distribution over the “typical” buyer’s value for item i .


Suppose there are m items and $\max_{i,v} f_i(v) \leq \kappa$.

Let $\hat{\boldsymbol{\rho}} \in \mathbb{R}^m$ be the price vector output by our DP algorithm.

With probability $1 - \gamma$,

$$U_{\mathcal{S}}(\hat{\boldsymbol{\rho}}) \geq \max U_{\mathcal{S}}(\boldsymbol{\rho}) - \tilde{O}\left(\frac{Hm}{|\mathcal{S}|\varepsilon} + \frac{\gamma}{|\mathcal{S}|\kappa\sqrt{m}}\right)$$

$U_{\mathcal{S}}(\boldsymbol{\rho}) = \text{average revenue}$



We also show that our DP algorithm has strong utility guarantees for many **mechanism design problems**.

Multiple bidders

Buyers with unit-demand valuations

Buyers with general valuations

2nd price auctions with reserve prices



Introduction

Setup

- Overview: Algorithm configuration and auction design

- Differential privacy

- Private pricing design

The algorithm

- Privacy guarantees

- Utility guarantees

- Example: private multi-item pricing

➡ Summary

We provide a **private** algorithm for maximizing data-dependent piecewise Lipschitz functions.

Utility function	Parameter Dimension	Privacy guarantee	Efficiency
Piecewise concave and Lipschitz	Multiple	(ϵ, δ)	✓
Piecewise Lipschitz	Single	$(\epsilon, 0)$	✓
Piecewise Lipschitz	Multiple	$(\epsilon, 0)$	✗

Questions?