

How much data is sufficient to learn high-performing algorithms?

Ellen Vitercik

UC Berkeley

Purdue University, Theory Seminar



Nina Balcan



Dan DeBlasio



Travis Dick



Carl Kingsford



Tuomas Sandholm



Appeared in STOC'21

Data-driven algorithm design

Algorithms often have **many tunable parameters**

Significant impact on runtime, solution quality, ...

Hand-tuning is **time-consuming**, **tedious**, and **error prone**





Data-driven algorithm design

Goal: Automate algorithm configuration via machine learning
Algorithmically find good parameter settings
using a set of "typical" inputs from application at hand

Training set

Parameter setting should – **ideally** – be good on future inputs



Example: Sequence alignment

Goal: Line up pairs of strings

Applications: Biology, natural language processing, etc.



Did you mean: [vitercik](#)

Sequence alignment algorithms

Input: Two sequences S and S'

Output: Alignment of S and S'

$S = A \ C \ T \ G$
 $S' = G \ T \ C \ A$

Gap
↓
A - - C T G
- G T C A -
↑ ↑ ↑
Insertion/deletion (*indel*) Match Mismatch


Sequence alignment algorithms

Standard algorithm with parameters $\rho_1, \rho_2, \rho_3 \geq 0$:

Return alignment maximizing:

$$(\# \text{ matches}) - \rho_1 \cdot (\# \text{ mismatches}) - \rho_2 \cdot (\# \text{ indels}) - \rho_3 \cdot (\# \text{ gaps})$$

$S = A \ C \ T \ G$
 $S' = G \ T \ C \ A$



Sequence alignment algorithms


Can sometimes access **ground-truth, reference** alignment

E.g., in computational biology: Bahr et al., Nucleic Acids Res.'01; Raghava et al., BMC Bioinformatics '03; Edgar, Nucleic Acids Res.'04; Walle et al., Bioinformatics'04

Requires extensive manual alignments
...rather just run parameterized algorithm

How to tune algorithm's parameters?

*"There is **considerable disagreement** among molecular biologists about the **correct choice**" [Gusfield et al. '94]*



A	-	-	C	T	G
-	G	T	C	A	-

Sequence alignment algorithms

-GRTCPKPDDL PFSTVVP-LKTFYEPGEEITYSCKPGYVSRGGMRKFICPLTGLWPINTLKCTP
E-VKCPFPSRPDNGFVNYPAPKPTLYYKDKATFGCHDGYSLDGP-EEIECTKLG NWSAMPSC-KA

Ground-truth alignment of protein sequences

Sequence alignment algorithms

-GRTCPKPDDL PFSTVVP-LKTFYEPGEEITYSCKPGYVSRGGMRKFICPLTGLWPINTLKCTP
E-VKCPFPSRPDNGFVNYPAPKPTLYYKDKATFGCHDGYSLDGP-EEIECTKLGNEWSAMPSC-KA

Ground-truth alignment of protein sequences

GRTCP---KPDDL PFSTVVPLKTFYEPGEEITYSCKPGYVSRGGMRKFICPLTGLWPINTLKCTP
EVKCPFPSRPDN-GFVNYPAPKPTLYYK-DKATFGCHDGY-SLDGP EEIECTKLGNEWS-AMPSCKA

Alignment by algorithm with **poorly-tuned** parameters

Sequence alignment algorithms

-GRTCPKPDDL PFSTVVP-LKTFYEPGEEITYSCKPGYVSRGGMRKFICPLTGLWPINTLKCTP
E-VKCPFPSRPDNGFVNYP AKPTLYYKDKATFGCHDGYSLDGP-EEIECTKLG NWSAMPSC-KA

Ground-truth alignment of protein sequences

GRTCP---KPDDL PFSTVVPLKTFYEPGEEITYSCKPGYVSRGGMRKFICPLTGLWPINTLKCTP
EVKCPFPSRPDN-GFVNYP AKPTLYYK-DKATFGCHDGY-SLDGPEEIECTKLG NWS-AMPSCKA

Alignment by algorithm with **poorly-tuned** parameters

GRTCPKPDDL PFSTV-VPLKTFYEPGEEITYSCKPGYVSRGGMRKFICPLTGLWPINTLKCTP
EVKCPFPSRPDNGFVNYP AKPTLYYKDKATFGCHDGY-SLDGPEEIECTKLG NWSA-MPSCKA

Alignment by algorithm with **well-tuned** parameters

Automated parameter tuning procedure

1. Fix parameterized algorithm
2. Receive training set T of "typical" inputs



3. Find parameters with good performance on average over T
Runtime, solution quality, etc.

Automated parameter tuning procedure

1. Fix parameterized algorithm
2. Receive training set T of "typical" inputs



3. Find parameters with good performance on average over T

Output alignment is close to reference alignment

Automated parameter tuning procedure

1. Fix parameterized algorithm
2. Receive training set T of "typical" inputs

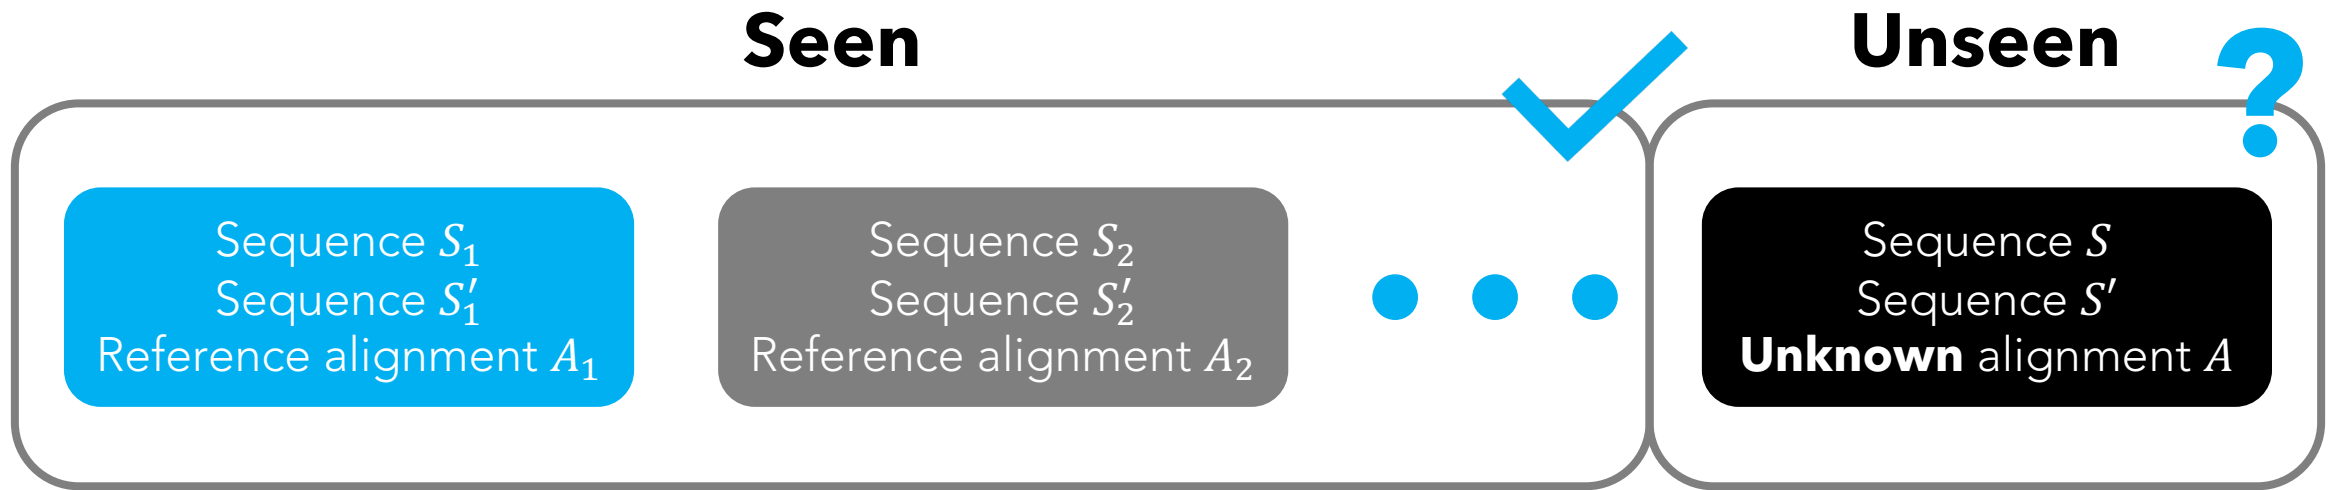


3. Find parameters with good performance on average over T

Key question (focus of talk):

Will those parameters have good **future** performance?

Automated parameter tuning procedure



Key question (focus of talk):

Will those parameters have good **future** performance?

Existing research



Constraint satisfaction

[Horvitz, Ruan, Gomes, Krautz, Selman, Chickering, UAI'01; ...]



Integer & linear programming

[Leyton-Brown, Nudelman, Andrew, McFadden, Shoham, CP '03; ...]



Economics (mechanism design)

[Likhodedov, Sandholm, AAAI '04, '05; ...]



Computational biology

[Majoros, Salzberg, Bioinformatics'04; ...]

**Applied
research**

2000

2021

Existing research

Automated algorithm configuration and selection

[Gupta, Roughgarden, ITCS'16; Balcan, Nagarajan, **Vitercik**, White, COLT'17; Balcan, Cambridge University Press '20; ...]

Learning-augmented algorithms

[Lykouris, Vassilvitskii, ICML'18; Mitzenmacher, NeurIPS'18; ...]

**Applied
research**

**Theory
research**

2000

2021

Theoretical guarantees are needed to build firm foundations

This talk: Main result

Key question (focus of talk):

Good performance on **average** over **training set** implies good **future** performance?

Answer this question for any parameterized algorithm where:

Performance is **piecewise-structured** function of parameters


Piecewise constant, linear, quadratic, ...

This talk: Main result


Performance is piecewise-structured function of parameters

Piecewise constant, linear, quadratic, ...

Algorithmic
performance
on fixed input



Piecewise constant




Piecewise linear



Piecewise ...


Example: Sequence alignment

Distance between **algorithm's output** given S, S'
and **ground-truth** alignment is p-wise constant




Piecewise structure

Piecewise structure unifies **seemingly disparate** problems:



Integer programming

Balcan, Dick, Sandholm, **V**, ICML'18
Balcan, Nagarajan, **V**, White, COLT'17




Computational biology

Balcan, DeBlasio, Dick, Kingsford,
Sandholm, **V**, STOC'21



Clustering

Balcan, Nagarajan, **V**, White, COLT'17
Balcan, Dick, White, NeurIPS'18
Balcan, Dick, Lang, ICLR'20



Greedy algorithms

Gupta, Roughgarden, ITCS'16



Mechanism configuration

Balcan, Sandholm, **V**, EC'18


Online configuration [Gupta, Roughgarden, ITCS'16, Cohen-Addad and Kanade, AISTATS'17]

Exploited piecewise-Lipschitz structure to provide regret bounds

[Balcan, Dick, **V**, FOCS'18; Balcan, Dick, Pegden, UAI'20; Balcan, Dick, Sharma, AISTATS'20]


Piecewise structure

Piecewise structure unifies **seemingly disparate** problems:




Integer programming

Balcan, Dick, Sandholm, **V**, ICML'18
Balcan, Nagarajan, **V**, White, COLT'17




Computational biology

Balcan, DeBlasio, Dick, Kingsford, Sandholm, **V**, STOC'21



Clustering

Balcan, Nagarajan, **V**, White, COLT'17
Balcan, Dick, White, NeurIPS'18
Balcan, Dick, Lang, ICLR'20



Greedy algorithms

Gupta, Roughgarden, ITCS'16



Mechanism configuration



Balcan, Sandholm, **V**, EC'18

Ties to a long line of research on machine learning for **revenue maximization**

Likhodedov, Sandholm, AAAI'04, '05; Balcan, Blum, Hartline, Mansour, FOCS'05; Elkind, SODA'07; Cole, Roughgarden, STOC'14; Mohri, Medina, ICML'14; Devanur, Huang, Psomas, STOC'16; ...

Primary challenge:

Algorithmic performance is a **volatile** function of parameters
Complex connection between parameters and performance



For well-understood functions in machine learning theory:
Simple connection between function parameters and value

Outline

1. Introduction
- 2. Model and problem formulation**
3. Our guarantees
4. Conclusion and future directions

Model

\mathbb{R}^d : Set of all parameters

\mathcal{X} : Set of all inputs

Example: Sequence alignment

\mathbb{R}^3 : Set of alignment algorithm parameters

\mathcal{X} : Set of sequence pairs



$S = A \ C \ T \ G$
 $S' = G \ T \ C \ A$

One sequence pair $x = (S, S') \in \mathcal{X}$

Algorithmic performance


$u_{\boldsymbol{\rho}}(x)$ = utility of algorithm parameterized by $\boldsymbol{\rho} \in \mathbb{R}^d$ on input x
E.g., runtime, solution quality, distance to ground truth, ...

Assume $u_{\boldsymbol{\rho}}(x) \in [-1, 1]$

Can be generalized to $u_{\boldsymbol{\rho}}(x) \in [-H, H]$

Algorithmic performance

$u_{\rho}(x)$ = distance between algorithm's output and ground-truth




One sequence pair $x = (S, S') \in \mathcal{X}$

Model

Standard assumption: Unknown distribution \mathcal{D} over inputs
Distribution models specific application domain at hand



E.g., distribution over pairs of DNA strands



E.g., distribution over pairs of protein sequences

Generalization bounds

Key question: For any parameter setting ρ ,
is **average** utility on training set close to **expected** utility?

Formally: Given samples $x_1, \dots, x_N \sim \mathcal{D}$, for any ρ ,

Generalization bounds

Key question: For any parameter setting ρ ,
is **average** utility on training set close to **expected** utility?

Formally: Given samples $x_1, \dots, x_N \sim \mathcal{D}$, for any ρ ,

$$\left| \underbrace{\frac{1}{N} \sum_{i=1}^N u_{\rho}(x_i)}_{\text{Empirical average utility}} - \mathbb{E}_{x \sim \mathcal{D}}[u_{\rho}(x)] \right| \leq ?$$

Generalization bounds

Key question: For any parameter setting ρ ,
is **average** utility on training set close to **expected** utility?

Formally: Given samples $x_1, \dots, x_N \sim \mathcal{D}$, for any ρ ,

$$\left| \frac{1}{N} \sum_{i=1}^N u_{\rho}(x_i) - \underbrace{\mathbb{E}_{x \sim \mathcal{D}}[u_{\rho}(x)]}_{\text{Expected utility}} \right| \leq ?$$

Generalization bounds

Key question: For any parameter setting ρ ,
is **average** utility on training set close to **expected** utility?

Formally: Given samples $x_1, \dots, x_N \sim \mathcal{D}$, for any ρ ,

$$\left| \frac{1}{N} \sum_{i=1}^N u_{\rho}(x_i) - \mathbb{E}_{x \sim \mathcal{D}}[u_{\rho}(x)] \right| \leq ?$$

Good **average empirical** utility  Good **expected** utility

Outline

1. Introduction
2. Model and problem formulation
3. Our guarantees
 - a. Example of piecewise-structured utility function**
 - b. Piecewise-structured functions more formally
 - c. Main theorem
 - d. Application: Sequence alignment
4. Conclusion and future directions

Sequence alignment algorithms

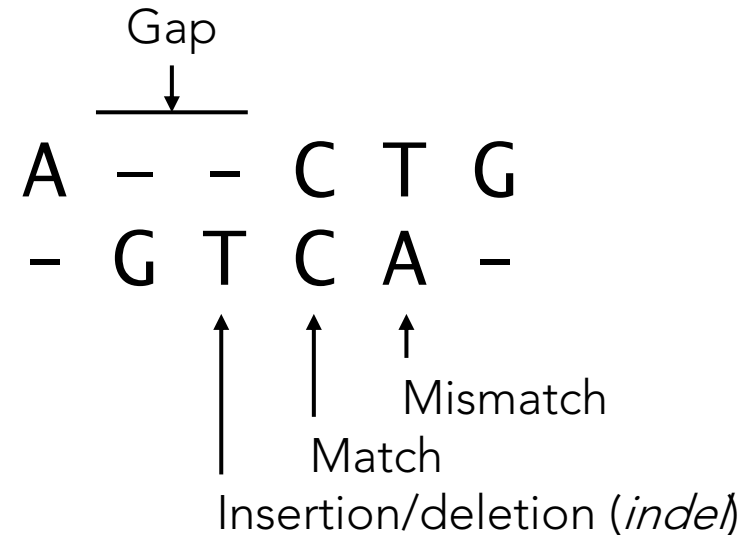
Standard algorithm with parameters $\rho_1, \rho_2, \rho_3 \geq 0$:

Return alignment maximizing:

$$(\# \text{ matches}) - \rho_1 \cdot (\# \text{ mismatches}) - \rho_2 \cdot (\# \text{ indels}) - \rho_3 \cdot (\# \text{ gaps})$$

$S = A \ C \ T \ G$

$S' = G \ T \ C \ A$




Sequence alignment algorithms

Lemma:

For any pair S, S' , there's a small partition of \mathbb{R}^3 s.t. in any region, algorithm's output is fixed across all parameters in region

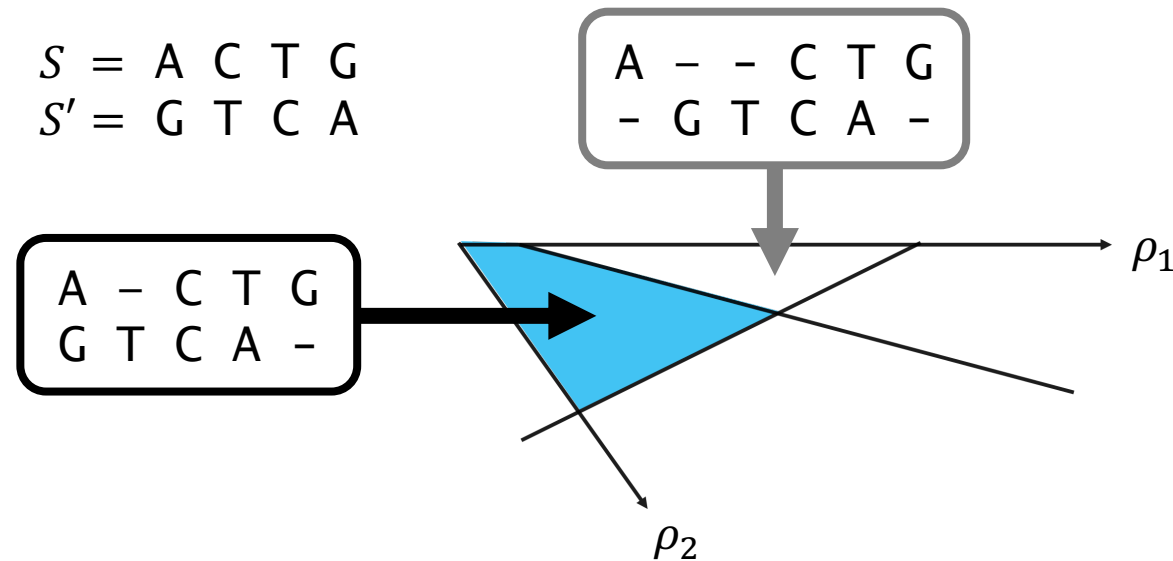
$S = A \ C \ T \ G$
 $S' = G \ T \ C \ A$



Sequence alignment algorithms

Lemma:

For any pair S, S' , there's a small partition of \mathbb{R}^3 s.t. in any region, algorithm's output is fixed across all parameters in region




Piecewise-constant utility function

Corollary:

Utility is piecewise constant function of parameters

Distance between algorithm's output and ground-truth alignment



Outline

1. Introduction
2. Model and problem formulation
3. Our guarantees
 - a. Example of piecewise-structured utility function
 - b. Piecewise-structured functions more formally**
 - c. Main theorem
 - d. Application: Sequence alignment
4. Conclusion and future directions

Primal & dual classes

$u_{\boldsymbol{\rho}}(x)$ = utility of algorithm parameterized by $\boldsymbol{\rho} \in \mathbb{R}^d$ on input x
 $\mathcal{U} = \{u_{\boldsymbol{\rho}}: \mathcal{X} \rightarrow \mathbb{R} \mid \boldsymbol{\rho} \in \mathbb{R}^d\}$ **"Primal" function class**

Typically, prove guarantees by bounding **complexity** of \mathcal{U}

VC dimension, pseudo-dimension, Rademacher complexity, ...

Primal & dual classes

$u_{\boldsymbol{\rho}}(x)$ = utility of algorithm parameterized by $\boldsymbol{\rho} \in \mathbb{R}^d$ on input x
 $\mathcal{U} = \{u_{\boldsymbol{\rho}}: \mathcal{X} \rightarrow \mathbb{R} \mid \boldsymbol{\rho} \in \mathbb{R}^d\}$ **"Primal" function class**

Typically, prove guarantees by bounding **complexity** of \mathcal{U}

Challenge: \mathcal{U} is gnarly

E.g., in sequence alignment:

- Each domain element is a pair of sequences
- Unclear how to plot or visualize functions $u_{\boldsymbol{\rho}}$
- No obvious notions of Lipschitz continuity or smoothness to rely on

Primal & dual classes

$u_{\boldsymbol{\rho}}(x)$ = utility of algorithm parameterized by $\boldsymbol{\rho} \in \mathbb{R}^d$ on input x
 $\mathcal{U} = \{u_{\boldsymbol{\rho}}: \mathcal{X} \rightarrow \mathbb{R} \mid \boldsymbol{\rho} \in \mathbb{R}^d\}$ **“Primal” function class**

$u_x^*(\boldsymbol{\rho})$ = utility as function of parameters


$$u_x^*(\boldsymbol{\rho}) = u_{\boldsymbol{\rho}}(x)$$

$\mathcal{U}^* = \{u_x^*: \mathbb{R}^d \rightarrow \mathbb{R} \mid x \in \mathcal{X}\}$ **“Dual” function class**


- Dual functions have simple, Euclidean domain
- Often have ample structure can use to bound complexity of \mathcal{U}

Piecewise-structured functions


Dual functions $u_x^*: \mathbb{R}^d \rightarrow \mathbb{R}$ are **piecewise-structured**




Clustering
algorithm
configuration




**Integer
programming**
algorithm
configuration



**Selling
mechanism**
configuration



Greedy
algorithm
configuration



**Computational
biology**
algorithm
configuration



**Voting
mechanism**
configuration

Piecewise-structured functions

Online algorithm configuration

Gupta, Roughgarden, ITCS'16, Cohen-Addad and Kanade, AISTATS'17

Problem instances arrive online, not necessarily i.i.d.

Exploited piecewise-Lipschitz structure to give regret bounds

Balcan, Dick, [V](#), FOCS'18; Balcan, Dick, Pegden, UAI'20; Balcan, Dick, Sharma, AISTATS'20

Regret bounds require additional structure:

Boundaries between pieces don't concentrate


Outline

1. Introduction
2. Model and problem formulation
3. Our guarantees
 - a. Example of piecewise-structured utility function
 - b. Piecewise-structured functions more formally
 - c. Main theorem**
 - d. Application: Sequence alignment
4. Conclusion and future directions

Intrinsic complexity


“Intrinsic complexity” of function class \mathcal{G}

- Measures how well functions in \mathcal{G} fit complex patterns
- Specific ways to quantify “intrinsic complexity”:
 - VC dimension
 - Pseudo-dimension



More complex


Less complex



VC dimension

Complexity measure for binary-valued function classes \mathcal{F}
(Classes of functions $f: \mathcal{Y} \rightarrow \{-1, 1\}$)

E.g., linear separators




VC dimension of \mathcal{F}

Size of the largest set $\mathcal{S} \subseteq \mathcal{Y}$

that can be labeled in all $2^{|\mathcal{S}|}$ ways by functions in \mathcal{F}

Example: \mathcal{F} = Linear separators in \mathbb{R}^2

$$\text{VCdim}(\mathcal{F}) \geq 3$$




VC dimension of \mathcal{F}

Size of the largest set $\mathcal{S} \subseteq \mathcal{Y}$

that can be labeled in all $2^{|\mathcal{S}|}$ ways by functions in \mathcal{F}

Example: \mathcal{F} = Linear separators in \mathbb{R}^2

$$\text{VCdim}(\mathcal{F}) \geq 3$$




$$\text{VCdim}(\mathcal{F}) \leq 3$$

$$\text{VCdim}(\{\text{Linear separators in } \mathbb{R}^d\}) = d + 1$$

Sample complexity using VC dimension

Theorem [Vapnik, Chervonenkis, '71]: For any dist. \mathcal{D} over \mathcal{Y} , given $N = \tilde{O}\left(\frac{\text{VCdim}(\mathcal{F})}{\epsilon^2}\right)$ samples $y_1, \dots, y_N \sim \mathcal{D}$, WHP $\forall f \in \mathcal{F}$,


$$\left| \frac{1}{N} \sum_{i=1}^N f(y_i) - \mathbb{E}_{y \sim \mathcal{D}}[f(y)] \right| \leq \epsilon$$



Pseudo-dimension

Complexity measure for real-valued function classes \mathcal{G}
(Classes of functions $g: \mathcal{Y} \rightarrow [0,1]$)





E.g., affine functions



Pseudo-dimension of \mathcal{G}

Size of the largest set $\{y_1, \dots, y_N\} \subseteq \mathcal{Y}$ s.t.:
for some *targets* $z_1, \dots, z_N \in \mathbb{R}$,
all 2^N above/below patterns achieved by functions in \mathcal{G}





Example: \mathcal{G} = Affine functions in \mathbb{R} $\text{Pdim}(\mathcal{G}) \geq 2$



Sample complexity using pseudo-dim

Theorem [Pollard, '84]: For any dist. \mathcal{D} over \mathcal{Y} ,
given $N = \tilde{O}\left(\frac{\text{Pdim}(\mathcal{G})}{\epsilon^2}\right)$ samples $y_1, \dots, y_N \sim \mathcal{D}$, WHP $\forall g \in \mathcal{G}$,


$$\left| \frac{1}{N} \sum_{i=1}^N g(y_i) - \mathbb{E}_{y \sim \mathcal{D}}[g(y)] \right| \leq \epsilon$$



Main result (informal)

Boundary functions $f_1, \dots, f_k \in \mathcal{F}$ partition \mathbb{R}^d s.t. in each region, $u_x^*(\boldsymbol{\rho}) = \textcolor{teal}{g}(\boldsymbol{\rho})$ for some $\textcolor{teal}{g} \in \mathcal{G}$.

Training set of size $\tilde{O}\left(\frac{1}{\epsilon^2}(\text{VCdim}(\mathcal{F}^*) + \text{Pdim}(\mathcal{G}^*)) \log k\right)$ implies
WHP $\forall \boldsymbol{\rho}, |\textcolor{teal}{avg} \text{ utility over training set} - \textcolor{teal}{exp} \text{ utility}| \leq \epsilon$



Main result (informal)

Boundary functions $f_1, \dots, f_k \in \mathcal{F}$ partition \mathbb{R}^d s.t. in each region, $u_x^*(\boldsymbol{\rho}) = g(\boldsymbol{\rho})$ for some $g \in \mathcal{G}$.

Training set of size $\tilde{O}\left(\frac{1}{\epsilon^2}(\text{VCdim}(\mathcal{F}^*) + \text{Pdim}(\mathcal{G}^*)) \log k\right)$ implies
WHP $\forall \boldsymbol{\rho}, |\text{avg utility over training set} - \text{exp utility}| \leq \epsilon$

\mathcal{F}, \mathcal{G} are typically very well structured

- \mathcal{G} = set of all **constant** functions $\Rightarrow \text{Pdim}(\mathcal{G}^*) = O(1)$
- \mathcal{G} = set of all **linear** functions in \mathbb{R}^d $\Rightarrow \text{Pdim}(\mathcal{G}^*) = O(d)$

Main result (informal)

Boundary functions $f_1, \dots, f_k \in \mathcal{F}$ partition \mathbb{R}^d s.t. in each region,
 $u_x^*(\boldsymbol{\rho}) = g(\boldsymbol{\rho})$ for some $g \in \mathcal{G}$.


Theorem:

$$\text{Pdim}(\mathcal{U}) = \tilde{O}\left((\text{VCdim}(\mathcal{F}^*) + \text{Pdim}(\mathcal{G}^*)) \log k\right)$$

↑
Primal function class $\mathcal{U} = \{u_{\boldsymbol{\rho}} \mid \boldsymbol{\rho} \in \mathbb{R}^d\}$

Key lemma


Each boundary function $f: \mathbb{R}^d \rightarrow \{-1, 1\}$ splits \mathbb{R}^d into 2 regions



Key lemma

Given D boundaries, how many sign patterns do they make?

$$\left| \left\{ \begin{pmatrix} f_1(\boldsymbol{\rho}) \\ \vdots \\ f_D(\boldsymbol{\rho}) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right| \leq ?$$



Key lemma

Given D boundaries, how many sign patterns do they make?

$$\left| \left\{ \begin{pmatrix} f_1(\boldsymbol{\rho}) \\ \vdots \\ f_D(\boldsymbol{\rho}) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right| \leq ?$$

Note: Sauer's lemma tells us that for any D points $\boldsymbol{\rho}_1, \dots, \boldsymbol{\rho}_D \in \mathbb{R}^d$

$$\left| \left\{ \begin{pmatrix} f(\boldsymbol{\rho}_1) \\ \vdots \\ f(\boldsymbol{\rho}_D) \end{pmatrix} : f \in \mathcal{F} \right\} \right| \leq (eD)^{\text{VCdim}(\mathcal{F})}$$

This is where transitioning to the dual comes in handy!

Key lemma

Given D boundaries, how many sign patterns do they make?

$$\left| \left\{ \begin{pmatrix} f_1(\boldsymbol{\rho}) \\ \vdots \\ f_D(\boldsymbol{\rho}) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right| \leq (eD)^{\text{VCdim}(\mathcal{F}^*)}$$

Proof ideas

For any problem instances x_1, \dots, x_N and targets $z_1, \dots, z_N \in \mathbb{R}$,


$$\left| \left\{ \begin{pmatrix} \text{sgn}(u_{\boldsymbol{\rho}}(x_1) - z_1) \\ \vdots \\ \text{sgn}(u_{\boldsymbol{\rho}}(x_N) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right| \leq ?$$

Switching to the dual functions,

$$\left| \left\{ \begin{pmatrix} \text{sgn}(u_{x_1}^*(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(u_{x_N}^*(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right| \leq ?$$

Proof ideas

$$\left| \left\{ \begin{pmatrix} \operatorname{sgn}(u_{x_1}^*(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \operatorname{sgn}(u_{x_N}^*(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right| \leq ?$$




Proof ideas

$$\left| \left\{ \begin{pmatrix} \text{sgn}(u_{x_1}^*(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(u_{x_N}^*(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right| \leq ?$$

The duals $u_{x_1}^*, \dots, u_{x_N}^*$ correspond to Nk boundary functions in \mathcal{F}


How many regions R_1, \dots, R_M in \mathbb{R}^d ? $M \leq (eNk)^{\text{VCdim}(\mathcal{F}^*)}$



Proof ideas

$$\left| \left\{ \begin{pmatrix} \text{sgn}(u_{x_1}^*(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(u_{x_N}^*(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in R_j \right\} \right| \leq ?$$

$\forall \boldsymbol{\rho} \in R_j$, duals are simultaneously structured: $u_{x_i}^*(\boldsymbol{\rho}) = g_i(\boldsymbol{\rho}), \forall i$



Proof ideas

$$\left| \left\{ \begin{pmatrix} \text{sgn}(u_{x_1}^*(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(u_{x_N}^*(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in R_j \right\} \right| \leq ?$$

$\forall \boldsymbol{\rho} \in R_j$, duals are simultaneously structured: $u_{x_i}^*(\boldsymbol{\rho}) = g_i(\boldsymbol{\rho}), \forall i$

$$\left| \left\{ \begin{pmatrix} \text{sgn}(g_1(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(g_N(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in R_j \right\} \right| \leq ?$$

Proof ideas

$$\left| \left\{ \begin{pmatrix} \text{sgn}(u_{x_1}^*(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(u_{x_N}^*(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in R_j \right\} \right| \leq ?$$

$\forall \boldsymbol{\rho} \in R_j$, duals are simultaneously structured: $u_{x_i}^*(\boldsymbol{\rho}) = g_i(\boldsymbol{\rho}), \forall i$

$$\left| \left\{ \begin{pmatrix} \text{sgn}(g_1(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(g_N(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in R_j \right\} \right| \leq \frac{(eN)^{\text{Pdim}(\mathcal{G}^*)}}{\quad}$$

Follows from key lemma

Proof ideas

$$\left| \left\{ \begin{pmatrix} \text{sgn}(u_{x_1}^*(\boldsymbol{\rho}) - z_1) \\ \vdots \\ \text{sgn}(u_{x_N}^*(\boldsymbol{\rho}) - z_N) \end{pmatrix} : \boldsymbol{\rho} \in \mathbb{R}^d \right\} \right|$$
$$\leq \underbrace{(eNk)^{\text{VCdim}(\mathcal{F}^*)}}_{\text{Number of regions}} \underbrace{(eN)^{\text{Pdim}(\mathcal{G}^*)}}_{\text{Number of sign patterns within each region}}$$

$\text{Pdim}(\mathcal{U})$ equals largest N s.t. $2^N \leq (eNk)^{\text{VCdim}(\mathcal{F}^*)} (eN)^{\text{Pdim}(\mathcal{G}^*)}$,
so $\text{Pdim}(\mathcal{U}) = \tilde{O}((\text{VCdim}(\mathcal{F}^*) + \text{Pdim}(\mathcal{G}^*)) \log k)$


Outline

1. Introduction
2. Model and problem formulation
3. Our guarantees
 - a. Example of piecewise-structured utility function
 - b. Piecewise-structured functions more formally
 - c. Main theorem
 - d. Application: Sequence alignment**
4. Conclusion and future directions

Piecewise constant dual functions


Lemma:

Utility is piecewise constant function of parameters



Sequence alignment guarantees

Theorem: Training set of size $\tilde{O}\left(\frac{\log(\text{seq. length})}{\epsilon^2}\right)$ implies WHP $\forall \boldsymbol{\rho}$,
|**avg** utility over training set - **exp** utility| $\leq \epsilon$



Outline

1. Introduction
2. Model and problem formulation
3. Our guarantees
- 4. Conclusion and future directions**


Conclusion

Guarantees for learning high-performing algorithm parameters

- Apply if performance is piecewise-structured function of parameters
- Proved by exploiting connections between primal and dual classes

Algorithm families from diverse domains exhibit this structure

- Clustering
- Economics (mechanism design)
- Integer programming
- Computational biology



Future research: Algorithms

This talk: Generalization guarantees

Apply to any configuration procedure (approximate, heuristic, optimal)

How to quickly find **provably** good parameters?

Growing body of work, but still many open questions

Kleinberg, Leyton-Brown, Lucier

IJCAI'17

Weisz, György, Szepesvári

ICML '18, '19

Kleinberg, Leyton-Brown, Lucier, Graham

NeurIPS'19

Weisz, György, Lin, Graham, Leyton-Brown, Szepesvári, Lucier


NeurIPS'20

Future research: Mapping to parameters

This talk: Learn **1** parameter setting that's good in expectation

Future research: More flexible approach


Learn **mapping** from inputs to parameter settings



Future research: Data-dependent bound



Strength of our results: **Input-distribution agnostic**
Apply to **any** distribution over instances



Can be loose for non-worst-case distributions

Guarantees that improve based on "niceness" of distribution?

- Covering-style analysis
- Rademacher complexity